

# NBA6921 Project Team8

Yuanhong Cao, Annabelle Gu, Yihui Guo

11/13/2021

```
rm(list=ls())
options(digits = 3, scipen = 999)
library(tidyverse)
library(ISLR)
library(jtools)
library(caret)
library(ROCR)
library(glmnet)
library(ggcorrplot)
library(cowplot)
library(lmtest)
library(corr)
library(dplyr)
library(ggplot2)
library(caret)
library(e1071)
library(MASS)
library(leaps)
library(randomForest)
library(rpart)
library(ranger)
library(gbm)
set.seed(2)
train <- read.csv("train.csv")
```

```
Score <- matrix(, nrow = 4, ncol = 1)
rownames(Score) <- c("Linear Regression", "Random Forest", "Elastic Net", "Boosting")
colnames(Score) <- c("RMSE")
```

Score

	RMSE
Linear Regression	NA
Random Forest	NA
Elastic Net	NA
Boosting	NA

```
train <- train[,!colnames(train) %in% c('Id', 'MSZoning', 'Street',
                                         'Alley', 'LotShape', 'LandContour',
                                         'Utilities', 'LotConfig', 'LandSlope',
                                         'Neighborhood', 'Condition1',
```

```
'Condition2', 'BldgType', 'HouseStyle',
'RoofStyle', 'RoofMatl', 'Exterior1st',
'Exterior2nd', 'MasVnrType',
'ExterQual', 'ExterCond', 'Foundation',
'BsmtQual', 'BsmtCond', 'BsmtExposure',
'BsmtFinType1', 'BsmtFinType2',
'Heating', 'HeatingQC', 'CentralAir',
'Electrical', 'KitchenQual',
'Functional', 'FireplaceQu',
'GarageType', 'GarageFinish',
'GarageQual', 'GarageCond',
'PavedDrive', 'PoolQC', 'Fence',
'MiscFeature', 'SaleType',
'SaleCondition', 'MasVnrArea']]
```

```
colSums(is.na(train))
```

MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond
0	259	0	0	0
YearBuilt	YearRemodAdd	BsmtFinSF1	BsmtFinSF2	BsmtUnfSF
0	0	0	0	0
TotalBsmtSF	X1stFlrSF	X2ndFlrSF	LowQualFinSF	GrLivArea
0	0	0	0	0
BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	BedroomAbvGr
0	0	0	0	0
KitchenAbvGr	TotRmsAbvGrd	Fireplaces	GarageYrBlt	GarageCars
0	0	0	81	0
GarageArea	WoodDeckSF	OpenPorchSF	EnclosedPorch	X3SsnPorch
0	0	0	0	0
ScreenPorch	PoolArea	MiscVal	MoSold	YrSold
0	0	0	0	0
SalePrice				
0				

```
train$LotFrontage[is.na(train$LotFrontage)] <- median(train$LotFrontage,
na.rm=TRUE)
train$GarageYrBlt[is.na(train$GarageYrBlt)] <- median(train$GarageYrBlt,
na.rm=TRUE)
```

```
colSums(is.na(train))
```

MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond
0	0	0	0	0
YearBuilt	YearRemodAdd	BsmtFinSF1	BsmtFinSF2	BsmtUnfSF
0	0	0	0	0
TotalBsmtSF	X1stFlrSF	X2ndFlrSF	LowQualFinSF	GrLivArea
0	0	0	0	0
BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	BedroomAbvGr
0	0	0	0	0
KitchenAbvGr	TotRmsAbvGrd	Fireplaces	GarageYrBlt	GarageCars
0	0	0	0	0
GarageArea	WoodDeckSF	OpenPorchSF	EnclosedPorch	X3SsnPorch

```

      0      0      0      0      0
ScreenPorch PoolArea MiscVal MoSold YrSold
      0      0      0      0      0
SalePrice
      0

```

```

train_ind <- sample(1:nrow(train), 4/5*nrow(train))
house_train <- train[train_ind,]
house_test <- train[-train_ind,]

```

## Variation

```
range(house_train$SalePrice, na.rm = TRUE)
```

```
[1] 35311 625000
```

```
quantile(house_train$SalePrice, na.rm = TRUE)
```

```

 0%    25%    50%    75%   100%
35311 130875 162700 213500 625000

```

```

quantile(house_train$SalePrice,
  probs = seq(from = 0, to = 1, by = .1),
  na.rm = TRUE)

```

```


 0%    10%    20%    30%    40%    50%    60%    70%    80%    90%   100%
35311 109000 125000 137000 147000 162700 179000 196950 229074 277150 625000

```

```

ggplot(data = house_train) +
  geom_histogram(mapping = aes(x = SalePrice), binwidth = 15000)

```



```
minprice <- min(house_test$SalePrice)
maxprice <- max(house_test$SalePrice)
```

```
house_train$SalePrice = log(house_train$SalePrice)
house_test$SalePrice = log(house_test$SalePrice)
train$SalePrice = log(train$SalePrice)
```

```
str(house_train)
```

```
'data.frame':  1168 obs. of  36 variables:
 $ MSSubClass   : int  70 20 20 20 60 60 60 20 50 20 ...
 $ LotFrontage  : int  60 69 70 73 71 92 75 72 60 70 ...
 $ LotArea      : int  11414 7162 10150 8899 12209 11764 9750 8872 10410 8400 ...
 $ OverallQual  : int  7 5 5 7 6 8 7 5 3 6 ...
 $ OverallCond  : int  8 7 5 5 5 7 6 8 4 3 ...
 $ YearBuilt    : int  1910 1966 1958 2007 2001 1999 1998 1965 1915 1957 ...
 $ YearRemodAdd : int  1993 1966 1958 2007 2002 2007 1998 2008 1950 1957 ...
 $ BsmtFinSF1   : int  0 0 456 24 690 524 975 595 0 189 ...
 $ BsmtFinSF2   : int  0 0 0 0 0 0 0 0 0 661 ...
 $ BsmtUnfSF    : int  728 876 456 1316 114 628 133 317 672 628 ...
 $ TotalBsmtSF  : int  728 876 912 1340 804 1152 1108 912 672 1478 ...
 $ X1stFlrSF    : int  1136 904 912 1340 804 1164 1108 912 694 1478 ...
 $ X2ndFlrSF    : int  883 0 0 0 1157 1106 989 0 520 0 ...
```

```

$ LowQualFinSF : int 0 0 0 0 0 0 0 0 0 0 ...
$ GrLivArea    : int 2019 904 912 1340 1961 2270 2097 912 1214 1478 ...
$ BsmtFullBath : int 0 0 0 0 1 0 1 1 0 1 ...
$ BsmtHalfBath : int 0 0 0 0 0 0 0 0 0 0 ...
$ FullBath     : int 1 1 1 2 2 2 2 1 1 1 ...
$ HalfBath     : int 0 0 0 0 1 1 1 0 0 1 ...
$ BedroomAbvGr : int 3 3 2 3 3 4 3 2 3 3 ...
$ KitchenAbvGr : int 1 1 1 1 1 1 1 1 1 1 ...
$ TotRmsAbvGrd : int 8 6 5 6 7 9 8 5 6 6 ...
$ Fireplaces   : int 0 0 0 0 1 1 1 0 0 2 ...
$ GarageYrBlt  : int 1997 1966 1958 2007 2001 1999 1998 1992 1998 1957 ...
$ GarageCars   : int 2 1 1 2 2 3 2 2 3 2 ...
$ GarageArea   : int 532 408 275 396 560 671 583 576 936 442 ...
$ WoodDeckSF   : int 509 0 0 100 125 132 253 0 216 114 ...
$ OpenPorchSF  : int 135 0 0 30 192 57 170 240 0 0 ...
$ EnclosedPorch : int 0 0 0 0 0 0 0 0 160 0 ...
$ X3SsnPorch   : int 0 0 0 0 0 0 0 0 0 0 ...
$ ScreenPorch  : int 0 0 0 0 0 0 0 0 0 216 ...
$ PoolArea     : int 0 0 0 0 0 0 0 0 0 0 ...
$ MiscVal      : int 0 0 0 0 0 0 0 0 0 0 ...
$ MoSold       : int 10 12 7 8 6 4 6 12 1 6 ...
$ YrSold       : int 2009 2008 2007 2007 2009 2010 2006 2008 2006 2009 ...
$ SalePrice    : num 12 11.6 11.6 12.1 12.3 ...

```

```

norm <- function(x) {
  (x - mean(x)) / sd(x)
}
denorm <- function(x,minval,maxval) {
  x*(maxval-minval) + minval
}
trainprice <- house_train$SalePrice
testprice <- house_test$SalePrice
totalprice <- train$SalePrice
train <- as.data.frame(lapply(train[,1:35], norm))
house_train <- as.data.frame(lapply(house_train[,1:35], norm))
house_test <- as.data.frame(lapply(house_test[,1:35], norm))

```

```

house_train["SalePrice"] <- trainprice
house_test["SalePrice"] <- testprice
train["SalePrice"] <- totalprice

```

```


num_cols = unlist(lapply(house_train, is.numeric))
# Create the correlation matrix
corr = cor(house_train[,num_cols])

```

```

ggcorrplot(corr,
  type = "full", lab = FALSE,
  legend.title = "Correlation Coefficient",
  colors = c("#053061", "white", "#67001f"),
  ggtheme = ggplot2::theme_void,
  outline.col = "white")

```

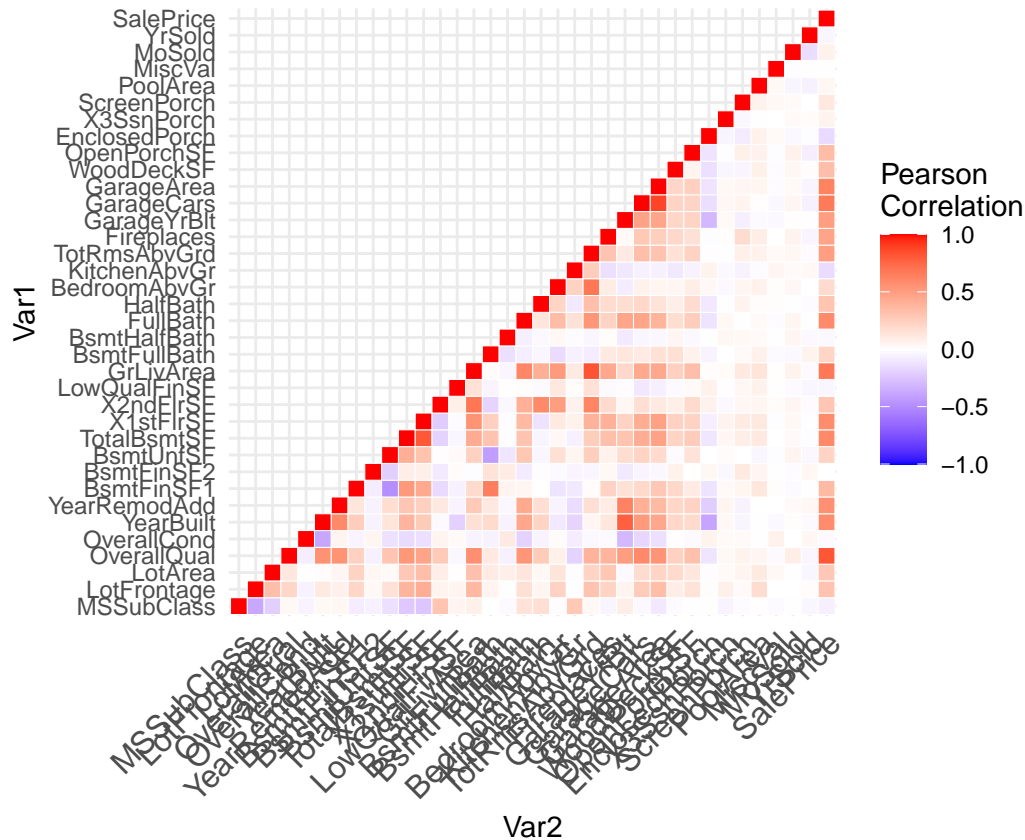


```
# Get lower triangle of the correlation matrix
get_lower_tri<-function(corr){
  corr[upper.tri(corr)] <- NA
  return(corr)
}

# Get upper triangle of the correlation matrix
get_upper_tri <- function(corr){
  corr[lower.tri(corr)]<- NA
  return(corr)
}

upper_tri <- get_upper_tri(corr)
# Melt the correlation matrix
library(reshape2)
melted_cormat <- melt(upper_tri, na.rm = TRUE)
# Heatmap
library(ggplot2)
ggplot(data = melted_cormat, aes(Var2, Var1, fill = value))+
  geom_tile(color = "white")+
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",
    midpoint = 0, limit = c(-1,1), space = "Lab",
    name="Pearson\nCorrelation") +
  theme_minimal()+
  theme(axis.text.x = element_text(angle = 45, vjust = 1,
    size = 12, hjust = 1))+
```


```
coord_fixed()
```



```
reorder_cormat <- function(corr){
  # Use correlation between variables as distance
  dd <- as.dist((1-corr)/2)
  hc <- hclust(dd)
  corr <- corr[hc$order, hc$order]
}

# Reorder the correlation matrix
corr <- reorder_cormat(corr)
upper_tri <- get_upper_tri(corr)
# Melt the correlation matrix
melted_cormat <- melt(upper_tri, na.rm = TRUE)
# Create a ggheatmap
ggheatmap <- ggplot(melted_cormat, aes(Var2, Var1, fill = value))+
  geom_tile(color = "white")+
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",
    midpoint = 0, limit = c(-1,1), space = "Lab",
    name="Pearson\nCorrelation") +
  theme_minimal()+ # minimal theme
  theme(axis.text.x = element_text(angle = 45, vjust = 1,
    size = 12, hjust = 1))+
  coord_fixed()
```

```
# Print the heatmap
print(ggheatmap)
```




```
#rank order for the cross-corr
library(lares)

corr_cross(house_train[,num_cols], # name of dataset
  max_pvalue = 0.05, # display only significant correlations (at 5% level)
  top = 10 # display top 10 couples of variables (by correlation coefficient)
)
```



## Ranked Cross-Correlations

10 most relevant




Correlations with p-value < 0.05

```
#focus on one variable vs the rest of all  
#looking at GarageCars since it appears to be the most correlated one  
#GarageCars is high correlated with GarageAreas therefore, we would adjust  
#those variables in our model  
corr_var(house_train[,num_cols], # name of dataset  
         GarageCars, # name of variable to focus on  
         top = 5 # display top 5 correlations  
)
```

## Correlations of GarageCars


Top 5 out of 11 variables (original & dummy)



```
corr_var(house_train[,num_cols], # name of dataset
  GrLivArea, # name of variable to focus on
  top = 5 # display top 5 correlations
)
```

## Correlations of GrLivArea

Top 5 out of 21 variables (original & dummy)



```
# Convert correlation matrix to data frame
corr_df = as_cordf(corr) %>%
# Focus on the Salary variable
  focus(SalePrice) %>%
# Get the absolute value of the correlation
# coefficient
  mutate(SalePrice = abs(SalePrice)) %>%
# Sort variables by absolute value of correlation
# coefficient
  arrange(SalePrice) %>%
# Clean up headers
  rename(`correlation with SalePrice` = term) %>%
  rename(corr_coef = SalePrice)
corr_df
```

```
# A tibble: 35 x 2
  `correlation with SalePrice` corr_coef
  <chr>                        <dbl>
1 BsmtFinSF2                  0.000198
2 MiscVal                    0.00540
3 OverallCond                 0.0264
4 BsmtHalfBath               0.0285
5 LowQualFinSF              0.0380
6 YrSold                    0.0385
7 PoolArea                  0.0418
8 X3SsnPorch                 0.0622
```

```

 9 MSSubClass                0.0679
10 MoSold                    0.0729
# ... with 25 more rows

```

```

# x = which(corr_df$corr_coef >= 0.5)
x = corr_df[which(corr_df$corr_coef >= 0.5),]
new_var = x['correlation with SalePrice']
# new_var
# house_train %>%
house_train <- house_train[,colnames(house_train) %in%
                           c(new_var$`correlation with SalePrice`,
                             'SalePrice')]
house_test <- house_test[,colnames(house_test) %in%
                          c(new_var$`correlation with SalePrice`,
                            'SalePrice')]
# house_train = house_train %>% select(new_var)
# house_test = house_test %>% select(new_var)

```

```
corr_df[which(corr_df$corr_coef >= 0.5),]
```

```

# A tibble: 11 x 2
  'correlation with SalePrice' corr_coef
  <chr>                      <dbl>
1 GarageYrBlt                0.501
2 TotRmsAbvGrd               0.502
3 YearRemodAdd               0.568
4 X1stFlrSF                  0.583
5 YearBuilt                  0.584
6 FullBath                   0.592
7 TotalBsmtSF                0.600
8 GarageArea                 0.634
9 GarageCars                 0.670
10 GrLivArea                  0.675
11 OverallQual                0.820

```

```




p1 <- ggplot(house_train,mapping = aes(x = TotRmsAbvGrd,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
             se=FALSE,colour = "blue")
p2 <- ggplot(house_train,mapping = aes(x =YearRemodAdd,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
             se=FALSE,colour = "blue")
p3 <- ggplot(house_train,mapping = aes(x =YearBuilt,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
             se=FALSE,colour = "blue")
p4 <- ggplot(house_train,mapping = aes(x = FullBath,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
             se=FALSE,colour = "blue")
p5 <- ggplot(house_train,mapping = aes(x =X1stFlrSF,y=SalePrice)) +
  geom_point() +

```


```

geom_smooth(method = "lm", formula = y~x,
            se=FALSE,colour = "blue")
p6 <- ggplot(house_train,mapping = aes(x =TotalBsmtSF,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
            se=FALSE,colour = "blue")
p7 <- ggplot(house_train,mapping = aes(x = GarageArea,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
            se=FALSE,colour = "blue")
p8 <- ggplot(house_train,mapping = aes(x =GarageCars,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
            se=FALSE,colour = "blue")
p9 <- ggplot(house_train,mapping = aes(x =GrLivArea,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
            se=FALSE,colour = "blue")
p10 <- ggplot(house_train,mapping = aes(x =OverallQual,y=SalePrice)) +
  geom_point() +
  geom_smooth(method = "lm", formula = y~x,
            se=FALSE,colour = "blue")
plot_grid(p1,p2,p3, ncol = 3)


```




```
plot_grid(p4,p5,p6, ncol = 3)
```



```
plot_grid(p7,p8,p9, ncol = 3)
```



```
plot_grid(p10, ncol = 3)
```



## Linear Regression

```
#build base model
lm1 <- lm(SalePrice~., data = house_train)
summary(lm1)
```

Call:

```
lm(formula = SalePrice ~ ., data = house_train)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.0151	-0.0750	0.0085	0.0906	0.5275

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	12.029148	0.004817	2497.19	< 0.0000000000000002	***
OverallQual	0.135660	0.007914	17.14	< 0.0000000000000002	***
YearBuilt	0.079051	0.008824	8.96	< 0.0000000000000002	***
YearRemodAdd	0.048187	0.006719	7.17	0.000000000000013	***
TotalBsmtSF	0.032870	0.008998	3.65	0.00027	***
X1stFlrSF	0.021118	0.009098	2.32	0.02044	*
GrLivArea	0.098481	0.010755	9.16	< 0.0000000000000002	***
FullBath	0.000301	0.007174	0.04	0.96653	



TotRmsAbvGrd	0.011784	0.008925	1.32	0.18698
GarageYrBlt	-0.034113	0.008367	-4.08	0.0000487672144 ***
GarageCars	0.049146	0.011088	4.43	0.0000101959037 ***
GarageArea	0.014874	0.010916	1.36	0.17328

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.165 on 1156 degrees of freedom

Multiple R-squared: 0.818, Adjusted R-squared: 0.817

F-statistic: 473 on 11 and 1156 DF, p-value: <0.0000000000000002

```
#check if there are any outliers
fit <- fitted(lm1)
stud.res <- studres(lm1)
stud.fit <- data.frame("fit"=fit,"stud.res"=stud.res)
ggplot(stud.fit, mapping = aes(x=fit,y=stud.res))+
geom_point()
```



```
#index1 <- which(stud.res > 5)
index2 <- which(stud.res < -5)
index <- index2
index
```

142 706 837

142 706 837

```
summary(lm1)$sigma
```

```
[1] 0.165
```

```
#summary(lm1)$r.squared
```

```
#remove outliers
```

```
adformula <- formula(SalePrice~.)  
lm_no_outlier = lm(adformula, data = house_train[-index,])  
summary(lm_no_outlier)$sigma
```

```
[1] 0.14
```

```
#summary(lm_no_outlier)$r.squared
```

```
#since FullBath & GarageArea & TotRmsAbvGrd 's pval is greater than 0.05,  
#we dont think it is statistically significant,we run again with a smaller model
```

```
lm2 <- lm(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd, data = house_train)  
summary(lm2)
```

Call:

```
lm(formula = SalePrice ~ . - FullBath - GarageArea - TotRmsAbvGrd,  
    data = house_train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.0103	-0.0761	0.0076	0.0912	0.5206

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	12.02915	0.00482	2496.80	< 0.0000000000000002 ***
OverallQual	0.13442	0.00788	17.06	< 0.0000000000000002 ***
YearBuilt	0.07711	0.00861	8.96	< 0.0000000000000002 ***
YearRemodAdd	0.04739	0.00668	7.09	0.000000000000023 ***
TotalBsmstSF	0.03298	0.00889	3.71	0.00022 ***
X1stFlrSF	0.02243	0.00905	2.48	0.01333 *
GrLivArea	0.10948	0.00680	16.10	< 0.0000000000000002 ***
GarageYrBlt	-0.03164	0.00813	-3.89	0.00011 ***
GarageCars	0.06172	0.00649	9.50	< 0.0000000000000002 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.165 on 1159 degrees of freedom

Multiple R-squared: 0.818, Adjusted R-squared: 0.817

F-statistic: 650 on 8 and 1159 DF, p-value: <0.0000000000000002

```
summary(lm2)$sigma
```

```
[1] 0.165
```

```
summary(lm2)$r.squared
```

```
[1] 0.818
```

```
#remove outliers
```

```
adformula <- formula(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd)
lm2_no_outlier = lm(adformula, data = house_train[-index,])
summary(lm2_no_outlier)$sigma
```

```
[1] 0.141
```

```
summary(lm2_no_outlier)$r.squared
```

```
[1] 0.865
```

```
#r^2 decreased to 0.761
```

```
#compare two linear regression model
```

```
anova(lm2_no_outlier,lm_no_outlier)
```

Analysis of Variance Table

```
Model 1: SalePrice ~ (OverallQual + YearBuilt + YearRemodAdd + TotalBsmtSF +
  X1stFlrSF + GrLivArea + FullBath + TotRmsAbvGrd + GarageYrBlt +
  GarageCars + GarageArea) - FullBath - GarageArea - TotRmsAbvGrd
Model 2: SalePrice ~ OverallQual + YearBuilt + YearRemodAdd + TotalBsmtSF +
  X1stFlrSF + GrLivArea + FullBath + TotRmsAbvGrd + GarageYrBlt +
  GarageCars + GarageArea
```

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	1156	22.9				
2	1153	22.4	3	0.499	8.54	0.000013 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
#Pval is smaller than 0.05, we chose lm2_no_outlier to perform our test
```

```
#outlier
```

```
house_train1 = house_train[-index,]
```

```
#best subset selection
```

```
# Draw validation set
```

```
house_validation_data = house_train1 %>% sample_frac(size = 0.3)
```

```
# Create the remaining training set
```

```
house_training_data = setdiff(house_train1, house_validation_data)
```

```
nvars = 7
```

```
regfit.best=regsubsets(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd,
  data=house_training_data,nvmax=nvars)
```

```
best.sum <- summary(regfit.best)
```

```
best.model <- which.max(best.sum$adjr2)
```

```
best.model
```

```
[1] 7
```

```
coef(regfit.best,id=best.model)
```

(Intercept)	OverallQual	YearBuilt	YearRemodAdd	TotalBsmtSF	GrLivArea
12.0337	0.1180	0.0795	0.0442	0.0812	0.1337
GarageYrBlt	GarageCars				
-0.0323	0.0593				

```
validation.mat=model.matrix(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd,
                             data=house_validation_data)
val.errors = numeric(nvars)

for(each in 1:nvars){
  coefi = coef(regfit.best,id=each)
  pred = validation.mat[,names(coefi)]%*%coefi
  val.errors[each]=mean((house_validation_data$SalePrice - pred)^2)
  sprintf("the val error is",val.errors[each])
}
best.subset.model = which.min(val.errors)
best.subset.model
```

```
[1] 7
```

```
#train on our test data in order to determine the accuracy
best.fit=regsubsets(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd,
                    data=house_train1,nvmax =7)
coefi_final1<- coef(best.fit,best.subset.model)
coefi_final1
```

(Intercept)	OverallQual	YearBuilt	YearRemodAdd	TotalBsmtSF	GrLivArea
12.0336	0.1153	0.0802	0.0457	0.0835	0.1431
GarageYrBlt	GarageCars				
-0.0316	0.0479				

```
#test data
test.mat1=model.matrix(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd,
                        data=house_test)

pred_test_lm = test.mat1[,names(coefi_final1)]%*%coefi_final1
```

```
head(house_test$SalePrice)
```

```
[1] 12.1 12.4 11.4 11.1 11.3 12.0
```

```
head(pred_test_lm)
```

```
 [,1]
1 12.0
2 12.5
```

```
3 11.6
4 11.3
5 11.3
6 12.0
```

```
pred_test_lm_org = exp(pred_test_lm)
```

```
library(Metrics)
#final lm model - accuracy
rmse(pred_test_lm,house_test$SalePrice)
```

```
[1] 0.184
```

```
#save as score
Score["Linear Regression","RMSE"] = rmse(pred_test_lm,house_test$SalePrice)

Score
```

	RMSE
Linear Regression	0.184
Random Forest	NA
Elastic Net	NA
Boosting	NA


## Random Forest

```
#build base model
rf <-randomForest(SalePrice ~., data=house_train)
rf
```

```
Call:
randomForest(formula = SalePrice ~ ., data = house_train)
Type of random forest: regression
Number of trees: 500
No. of variables tried at each split: 3
```

```
Mean of squared residuals: 0.0221
% Var explained: 85
```

```
plot(rf)
```



```
# number of trees with lowest MSE
which.min(rf$mse)
```

```
[1] 280
```

```
# RMSE of this optimal random forest
sqrt(rf$mse[which.min(rf$mse)])
```

```
[1] 0.148
```

```
#tuning parameter
house_tree_tune <- rpart(SalePrice ~ .,data = house_train,method="anova",
                          maxdepth=7)
house_tree_tune
```

```
n= 1168
```

```
node), split, n, deviance, yval
      * denotes terminal node
```

- 1) root 1168 172.00 12.0
- 2) OverallQual< 0.288 728 53.70 11.8
- 4) GrLivArea< -0.276 439 26.70 11.7
- 8) YearBuilt< -0.659 139 9.60 11.5

```

16) OverallQual< -1.9 16 1.45 11.1 *
17) OverallQual>=-1.9 123 5.22 11.6 *
9) YearBuilt>=-0.659 300 9.34 11.8
18) TotalBsmtSF< -0.112 163 4.45 11.7 *
19) TotalBsmtSF>=-0.112 137 2.57 11.9 *
5) GrLivArea>=-0.276 289 15.30 12.0
10) OverallQual< -0.441 120 5.59 11.9 *
11) OverallQual>=-0.441 169 6.38 12.1 *
3) OverallQual>=0.288 440 38.80 12.4
6) OverallQual< 1.02 259 10.50 12.2
12) GrLivArea< 1.1 221 6.76 12.2 *
13) GrLivArea>=1.1 38 1.07 12.5 *
7) OverallQual>=1.02 181 14.20 12.6
14) OverallQual< 1.75 133 7.20 12.5 *
15) OverallQual>=1.75 48 3.24 12.8 *

```

```

# hyperparameter grid search
hyper_grid <- expand.grid(
  mtry = seq(5, 10, by = 1),
  node_size = seq(4, 16, by = 2),
  sample_size = c(.5, .6, .70, .80),
  OOB_RMSE = 0
)
# total number of combinations
nrow(hyper_grid)

```

```
[1] 168
```

```

for(i in 1:nrow(hyper_grid)) {
  # train model
  model <- ranger(
    formula = SalePrice ~ .,
    data = house_train,
    num.trees = 348,
    mtry = hyper_grid$mtry[i],
    min.node.size = hyper_grid$node_size[i],
    sample.fraction = hyper_grid$sample_size[i] )
  # add OOB error to grid
  hyper_grid$OOB_RMSE[i] <- sqrt(model$prediction.error)
}

```

```

hyper_grid %>%
  arrange(OOB_RMSE) %>% head(10)

```

	mtry	node_size	sample_size	OOB_RMSE
1	6	6	0.7	0.147
2	7	6	0.8	0.147
3	6	8	0.8	0.147
4	5	12	0.8	0.147
5	6	4	0.6	0.147
6	8	14	0.8	0.147
7	6	12	0.7	0.147

8	6	4	0.8	0.147
9	5	6	0.7	0.147
10	5	6	0.8	0.147

```
best.rf <- hyper_grid %>%
  arrange(OOB_RMSE) %>%
  head(1)
best.rf
```

	mtry	node_size	sample_size	OOB_RMSE
1	6	6	0.7	0.147

```
optimal_rf <- ranger(
  formula = SalePrice ~ .,
  data = house_train,
  num.trees = 348,
  mtry = best.rf$mtry,
  min.node.size = best.rf$node_size,
  sample.fraction = best.rf$sample_size,
  importance = 'impurity')
```

```
#make predictions on
predict_rf <- predict(optimal_rf, house_test)$predictions
```

```
#store them in Score
Score["Random Forest", "RMSE"] = RMSE(predict_rf, house_test$SalePrice)

Score
```

	RMSE
Linear Regression	0.184
Random Forest	0.192
Elastic Net	NA
Boosting	NA

## Elastic net

```
# Predictor variables
x <- model.matrix(SalePrice~., house_train)[,-1]
# Outcome variable
y <- house_train$SalePrice

# Build the model using the training set
set.seed(123)
model <- train(
  SalePrice ~., data = house_train, method = "glmnet",
  trControl = trainControl("cv", number = 10),
  tuneLength = 10
)
# Best tuning parameter
model$bestTune
```



```
alpha lambda
5 0.1 0.00957
```

```
coef(model$finalModel, model$bestTune$lambda)
```

```
12 x 1 sparse Matrix of class "dgCMatrix"
```

```
      s1
(Intercept) 12.0291
OverallQual  0.1327
YearBuilt    0.0707
YearRemodAdd 0.0462
TotalBsmtSF  0.0337
X1stFlrSF    0.0222
GrLivArea    0.0924
FullBath     0.0024
TotRmsAbvGrd 0.0144
GarageYrBlt  -0.0240
GarageCars   0.0485
GarageArea   0.0157
```

```
# Make predictions on the test data
x.test <- model.matrix(SalePrice ~., house_test)[,-1]
predictions <- model %>% predict(x.test)
```

```
Score["Elastic Net","RMSE"] = RMSE(predictions, house_test$SalePrice)
```

```
Score
```

	RMSE
Linear Regression	0.184
Random Forest	0.192
Elastic Net	0.191
Boosting	NA


## Boosting

- Basic GBM model

```
hit_gbm <- gbm(
  formula = SalePrice ~ .,
  data = house_train,
  distribution = "gaussian", # SSE loss function
  n.trees = 1000,
  shrinkage = 0.001, #learning rate
  cv.folds = 10,
  interaction.depth = 5 #depth of each tree
)
# find index for number trees with minimum CV error
best <- which.min(hit_gbm$cv.error)
# get MSE and compute RMSE
sqrt(hit_gbm$cv.error[best])
```

```
[1] 0.229
```

```
gbm.perf(hit_gbm, method = "cv")
```




```
[1] 1000
```

```
hit_gbm <- gbm(  
  formula = SalePrice ~ .,  
  data = house_train,  
  distribution = "gaussian", # SSE loss function  
  n.trees = 10000,  
  shrinkage = 0.001, #learning rate  
  cv.folds = 10,  
  interaction.depth = 5 #depth of each tree  
)  
# find index for number trees with minimum CV error  
best <- which.min(hit_gbm$cv.error)  
# get MSE and compute RMSE  
sqrt(hit_gbm$cv.error[best])
```

```
[1] 0.148
```

```
gbm.perf(hit_gbm, method = "cv")
```



```
[1] 8176
```

```
pred.gbm.final <- predict.gbm(hit_gbm, n.trees=4000, newdata = house_test)
rmse.gbm.final.rmse <- sqrt(mean((house_test$SalePrice -
                                pred.gbm.final)^2))
rmse.gbm.final.rmse
```

```
[1] 0.188
```

```
pred.gbm.final
```

```
[1] 12.0 12.6 11.7 11.2 11.5 11.9 12.1 11.5 12.3 11.5 11.8 12.2 12.1 12.6 12.1
[16] 12.2 12.0 12.3 11.7 12.5 12.1 12.0 11.6 11.8 11.9 12.1 11.7 12.6 12.1 11.7
[31] 12.2 12.2 12.0 11.9 11.7 11.8 12.4 11.9 12.2 12.2 12.2 11.7 11.8 12.3 12.8
[46] 11.8 12.4 12.4 12.4 12.1 12.0 11.6 12.5 12.1 11.8 12.5 11.5 11.6 11.7 12.3
[61] 12.5 12.8 12.6 11.7 11.4 12.0 12.3 12.2 11.6 11.7 12.0 12.4 11.6 11.7 12.8
[76] 12.3 11.6 11.4 12.2 12.0 11.9 12.5 11.9 11.9 11.8 12.2 12.2 12.7 12.2 11.6
[91] 12.9 12.2 11.4 11.7 11.7 11.9 11.6 11.9 11.5 11.3 12.7 12.4 11.9 11.8 11.5
[106] 11.8 12.2 11.6 12.7 12.0 12.5 11.9 12.5 12.4 11.8 12.0 11.8 12.2 11.7 11.8
[121] 12.2 12.4 12.5 12.8 11.9 12.7 11.9 12.4 12.2 12.4 12.3 11.4 12.2 11.6 12.2
[136] 12.0 12.0 12.2 12.8 11.8 12.4 12.9 11.6 11.2 12.6 11.9 11.8 11.7 12.4 12.4
```

```
[151] 11.8 12.1 12.3 11.8 12.7 12.1 11.8 11.7 12.2 11.6 11.9 11.9 11.6 12.4 11.9
[166] 11.7 12.4 12.0 12.3 11.9 12.2 11.7 11.8 11.9 12.3 11.7 12.1 11.9 11.8 12.1
[181] 11.8 12.2 12.0 12.1 12.1 12.6 12.5 12.3 12.1 11.9 11.7 11.7 11.8 11.7 12.7
[196] 12.1 11.0 11.5 12.3 12.1 12.7 12.0 11.4 12.4 11.8 11.7 11.9 11.7 12.6 11.5
[211] 12.1 11.9 12.1 11.8 11.8 11.5 12.5 11.8 12.1 11.6 12.2 11.8 11.8 11.7 12.1
[226] 12.3 12.3 11.8 11.7 12.5 11.8 12.1 11.6 12.0 12.8 11.5 11.7 11.7 12.2 12.2
[241] 12.5 11.9 12.1 12.9 12.0 11.7 12.2 11.8 12.5 11.8 11.9 12.1 11.8 12.0 12.1
[256] 12.7 11.9 12.0 12.2 12.9 11.9 11.9 11.9 11.9 11.9 11.8 12.4 12.2 12.2 11.4
[271] 11.3 11.8 11.5 11.9 11.3 11.7 12.0 12.6 12.6 12.3 12.0 11.5 12.2 11.8 12.3
[286] 12.6 11.9 11.5 12.7 12.5 12.3 11.9
```

```
CV_RSq <- (cor(pred.gbm.final, house_test$SalePrice))^2
CV_RSq
```

```
[1] 0.843
```

```
# create hyperparameter grid
hyper_grid <- expand.grid(
  shrinkage = c(.001, .1),
  interaction.depth = c(1, 5),
  n.minobsinnode = c(5, 10),
  bag.fraction = c(.7, .8),
  optimal_trees = 0,
  min_RMSE = 0
)

# total number of combinations
nrow(hyper_grid)
```

```
[1] 16
```

```
# grid search
for(i in 1:nrow(hyper_grid)) {
  print(i)
  # train model
  gbm.tune <- gbm(
    formula = SalePrice ~ .,
    distribution = "gaussian",
    data = house_train,
    n.trees = 4000,
    interaction.depth = hyper_grid$interaction.depth[i],
    shrinkage = hyper_grid$shrinkage[i],
    n.minobsinnode = hyper_grid$n.minobsinnode[i],
    bag.fraction = hyper_grid$bag.fraction[i],
    cv.folds = 10)

  # add min training error and trees to grid
  hyper_grid$optimal_trees[i] <- which.min(gbm.tune$cv.error)
  hyper_grid$min_RMSE[i] <- sqrt(min(gbm.tune$cv.error))
}
```

```
[1] 1
```

```
[1] 2
[1] 3
[1] 4
[1] 5
[1] 6
[1] 7
[1] 8
[1] 9
[1] 10
[1] 11
[1] 12
[1] 13
[1] 14
[1] 15
[1] 16
```

---

```
hyper_grid %>%
  arrange(min_RMSE) %>%
  head(10)
```

	shrinkage	interaction.depth	n.minobsinnode	bag.fraction	optimal_trees
1	0.100	5	5	0.8	96
2	0.100	5	5	0.7	111
3	0.100	5	10	0.8	98
4	0.100	5	10	0.7	150
5	0.001	5	10	0.7	4000
6	0.001	5	5	0.7	4000
7	0.001	5	5	0.8	4000
8	0.001	5	10	0.8	4000
9	0.100	1	10	0.7	293
10	0.100	1	10	0.8	468

	min_RMSE
1	0.147
2	0.147
3	0.148
4	0.150
5	0.150
6	0.150
7	0.150
8	0.151
9	0.151
10	0.152

```
best.model <- hyper_grid %>%
  arrange(min_RMSE) %>%
  head(1)
best.model
```

	shrinkage	interaction.depth	n.minobsinnode	bag.fraction	optimal_trees
1	0.1	5	5	0.8	96

	min_RMSE
1	0.147

- 
- Let's re-run the GBM model with optimal hyper parameters

```
hit_gbm.final <- gbm(
  formula = SalePrice ~ .,
  data = house_train,
  distribution = "gaussian",
  n.trees = 4000,
  interaction.depth = best.model$interaction.depth,
  shrinkage = best.model$shrinkage,
  n.minobsinnode = best.model$n.minobsinnode,
  bag.fraction = best.model$bag.fraction,
  cv.folds = 10)
# find index for number trees with minimum CV error
best <- which.min(hit_gbm.final$cv.error)
# get MSE and compute RMSE
sqrt(hit_gbm.final$cv.error[best])
```

```
[1] 0.149
```

- 
- Make predictions on the test data

```
pred.gbm.final <- predict.gbm(hit_gbm.final, n.trees=4000, newdata = house_test)
rmse.gbm.final.rmse <- sqrt(mean((house_test$SalePrice -
  pred.gbm.final)^2))
rmse.gbm.final.rmse
```

```
[1] 0.205
```

```
pred.gbm.final
```

```
[1] 12.1 12.6 11.7 11.2 11.4 12.0 12.1 11.6 12.3 11.7 11.7 12.2 12.1 12.6 12.1
[16] 12.2 12.0 12.3 11.7 12.4 12.1 12.1 11.6 11.9 12.0 12.1 11.8 12.6 12.0 11.7
[31] 12.2 12.3 11.9 11.9 11.6 11.9 12.3 12.0 12.1 12.3 12.1 11.7 11.8 12.3 12.8
[46] 11.9 12.4 12.3 12.3 12.2 12.0 11.7 12.4 12.1 11.9 12.5 11.4 11.6 11.7 12.2
[61] 12.5 12.8 12.6 11.6 11.0 12.1 12.3 12.2 11.8 11.7 12.1 12.4 11.8 11.7 12.9
[76] 12.2 11.5 11.5 12.2 12.0 12.0 12.6 11.9 11.9 11.9 12.2 12.2 12.7 12.3 11.8
[91] 13.1 12.2 11.4 11.7 11.7 11.8 11.6 11.9 11.5 11.4 12.4 12.4 11.9 11.8 11.5
[106] 11.8 12.2 11.5 12.7 11.9 12.4 12.0 12.6 12.5 11.9 11.9 11.7 12.3 11.6 11.7
[121] 12.2 12.4 12.5 12.7 11.9 12.6 11.9 12.4 12.2 12.4 12.3 11.5 12.5 11.7 12.4
[136] 12.0 12.1 12.4 12.8 11.8 12.3 12.9 11.6 11.0 12.6 11.9 11.8 11.6 12.5 12.4
[151] 11.8 12.0 12.4 11.7 12.7 12.1 11.8 11.7 12.1 11.4 12.0 11.9 11.7 12.3 11.9
[166] 11.7 12.4 11.8 12.3 11.5 12.2 11.7 11.9 12.0 12.2 11.6 12.1 11.7 11.8 12.1
[181] 11.8 12.2 12.0 12.1 12.0 12.7 12.6 12.3 12.1 11.9 11.7 11.7 11.7 11.6 12.7
[196] 12.1 10.8 11.5 12.3 12.1 12.8 12.0 11.6 12.5 11.9 11.8 11.9 11.8 12.7 11.5
[211] 12.1 11.9 12.1 11.7 12.0 11.4 12.6 11.8 12.1 11.6 12.2 11.9 11.9 11.5 12.1
[226] 12.3 12.4 11.8 11.6 12.5 11.8 12.2 11.6 12.0 12.8 11.4 11.7 11.6 12.2 12.3
[241] 12.5 12.0 12.0 13.1 12.0 11.8 12.2 11.9 12.4 11.8 11.9 12.1 11.7 11.9 12.1
[256] 12.8 11.9 12.1 12.2 13.0 11.9 12.0 12.0 11.8 11.9 11.9 12.5 12.1 12.3 11.4
[271] 11.1 11.8 11.4 11.8 11.5 11.7 12.0 12.6 12.6 12.3 11.9 11.6 12.3 11.7 12.4
[286] 12.7 12.0 11.6 12.8 12.5 12.1 11.8
```

```
#store them in Score
Score["Boosting","RMSE"] = RMSE(pred.gbm.final, house_test$SalePrice)
Score
```

	RMSE
Linear Regression	0.184
Random Forest	0.192
Elastic Net	0.191
Boosting	0.205

```
test <- read.csv("test.csv")
test <- test[,!colnames(test) %in% c('Id','MSZoning','Street','Alley',
    'LotShape', 'LandContour', 'Utilities',
    'LotConfig', 'LandSlope', 'Neighborhood',
    'Condition1', 'Condition2', 'BldgType',
    'HouseStyle', 'RoofStyle', 'RoofMatl',
    'Exterior1st', 'Exterior2nd', 'MasVnrType',
    'ExterQual', 'ExterCond', 'Foundation',
    'BsmtQual', 'BsmtCond', 'BsmtExposure',
    'BsmtFinType1', 'BsmtFinType2', 'Heating',
    'HeatingQC', 'CentralAir', 'Electrical',
    'KitchenQual', 'Functional', 'FireplaceQu',
    'GarageType', 'GarageFinish', 'GarageQual',
    'GarageCond', 'PavedDrive', 'PoolQC',
    'Fence', 'MiscFeature', 'SaleType',
    'SaleCondition', 'MasVnrArea')]
```

```
final.fit=regsubsets(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd,
    data=train[,-index],nvmax =7)
```

Reordering variables and trying again:

```
final_model <- coef(final.fit,best.subset.model)
final_model
```

(Intercept)	MSSubClass	OverallQual	OverallCond	YearBuilt	BsmtHalfBath
12.02405	-0.04467	0.26503	0.04150	0.08609	0.00795
HalfBath	OpenPorchSF				
0.03783	0.02388				

```
#test data
final_test=model.matrix(SalePrice~.-FullBath-GarageArea-TotRmsAbvGrd,
    data=train[,-index])

final_test_lm = final_test[,names(final_model)]%*%final_model
```